

JupyterLab interface showing the first cell of code execution. The browser address bar shows `localhost:8888/notebooks/Red_Wine_Prediction/Wine_prediction.ipynb`. The JupyterLab title bar indicates the last checkpoint was 26 minutes ago. The code cell contains the following Python code:

```
[17]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
sns.set()
from warnings import filterwarnings
filterwarnings(action='ignore')

[18]: wine = pd.read_csv("winequality-red.csv")
print("Successfully Imported Data!")
wine.head()
```

The output shows the first five rows of the dataset:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	5
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	9.8	5
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	9.8	5
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	9.8	6
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	5

JupyterLab interface showing the second cell of code execution. The browser address bar shows `localhost:8888/notebooks/Red_Wine_Prediction/Wine_prediction.ipynb`. The JupyterLab title bar indicates the last checkpoint was 28 minutes ago. The code cell contains the following Python code:

```
[19]: print(wine.shape)

(1599, 12)

[20]: wine.describe(include='all')
```

The output shows the shape of the dataset and a detailed description of all columns:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
count	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000
mean	8.319637	0.527821	0.270976	2.538806	0.087467	15.874922	46.467792	0.996747	3.311113	0.658149	10.422983	5.636023
std	1.741096	0.179060	0.194801	1.409928	0.047065	10.460157	32.895324	0.001887	0.154386	0.169507	1.065668	0.807569
min	4.600000	0.120000	0.000000	0.900000	0.012000	1.000000	6.000000	0.990070	2.740000	0.330000	8.400000	3.000000
25%	7.100000	0.390000	0.090000	1.900000	0.070000	7.000000	22.000000	0.995600	3.210000	0.550000	9.500000	5.000000
50%	7.900000	0.520000	0.260000	2.200000	0.079000	14.000000	38.000000	0.996750	3.310000	0.620000	10.200000	6.000000
75%	9.200000	0.640000	0.420000	2.600000	0.090000	21.000000	62.000000	0.997835	3.400000	0.730000	11.100000	6.000000
max	15.900000	1.580000	1.000000	15.500000	0.611000	72.000000	289.000000	1.003690	4.010000	2.000000	14.900000	8.000000

The final line of code and its output are:

```
[21]: print(wine.isna().sum())

fixed acidity      0
volatile acidity   0
citric acid        0
```

JupyterLab interface showing a notebook titled "Wine\_prediction". The code cell [22] contains the following Python code:

```
residual sugar
chlorides
free sulfur dioxide
total sulfur dioxide
density
pH
sulphates
alcohol
quality
dtype: int64

[22]: wine.corr()
```

The output of the code is a correlation matrix for the wine quality dataset. The matrix shows the relationship between various chemical and physical properties and the quality of the wine. The diagonal elements are all 1.000000, indicating perfect self-correlation. The off-diagonal elements represent the correlation coefficients between different features.

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
fixed acidity	1.000000	-0.256131	0.671703	0.114777	0.093705	-0.153794	-0.113181	0.668047	-0.682978	0.183006	-0.061668	0.124052
volatile acidity	-0.256131	1.000000	-0.552496	0.001918	0.061298	-0.010504	0.076470	0.022026	0.234937	-0.260987	-0.202288	-0.390558
citric acid	0.671703	-0.552496	1.000000	0.143577	0.203823	-0.060978	0.035533	0.364947	-0.541904	0.312770	0.109903	0.226373
residual sugar	0.114777	0.001918	0.143577	1.000000	0.055610	0.187049	0.203028	0.355283	-0.085652	0.005527	0.042075	0.013732
chlorides	0.093705	0.061298	0.203823	0.055610	1.000000	0.005562	0.047400	0.200632	-0.265026	0.371260	-0.221141	-0.128907
free sulfur dioxide	-0.153794	-0.010504	-0.060978	0.187049	0.005562	1.000000	0.667666	-0.021946	0.070377	0.051658	-0.069408	-0.050656


JupyterLab interface showing a notebook titled "Wine\_prediction". The code cell [23] contains the following Python code:

```
[23]: wine.groupby('quality').mean()
```


The output of the code is a grouped mean table for the wine quality dataset. The table shows the mean values of various chemical and physical properties for each quality level (3, 4, 5, 6, 7, 8). The columns are ordered alphabetically by feature name.

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
quality											
3	8.360000	0.884500	0.171000	2.635000	0.122500	11.000000	24.900000	0.997464	3.398000	0.570000	9.955000
4	7.779245	0.693962	0.174151	2.694340	0.090679	12.264151	36.245283	0.996542	3.381509	0.596415	10.265094
5	8.167254	0.577041	0.243686	2.528855	0.092736	16.983847	56.513950	0.997104	3.304949	0.620969	9.899706
6	8.347179	0.497484	0.273824	2.477194	0.084956	15.711599	40.869906	0.996615	3.318072	0.675329	10.629519
7	8.872362	0.403920	0.375176	2.720603	0.076588	14.045226	35.020101	0.996104	3.290754	0.741256	11.465913
8	8.566667	0.423333	0.391111	2.577778	0.068444	13.277778	33.444444	0.995212	3.267222	0.767778	12.094444


```
sns.countplot(wine['quality'])  
plt.show()
```




```
sns.countplot(wine['pH'])  
plt.show()
```




```
sns.countplot(wine['alcohol'])  
plt.show()
```




```
sns.countplot(wine['fixed acidity'])  
plt.show()
```



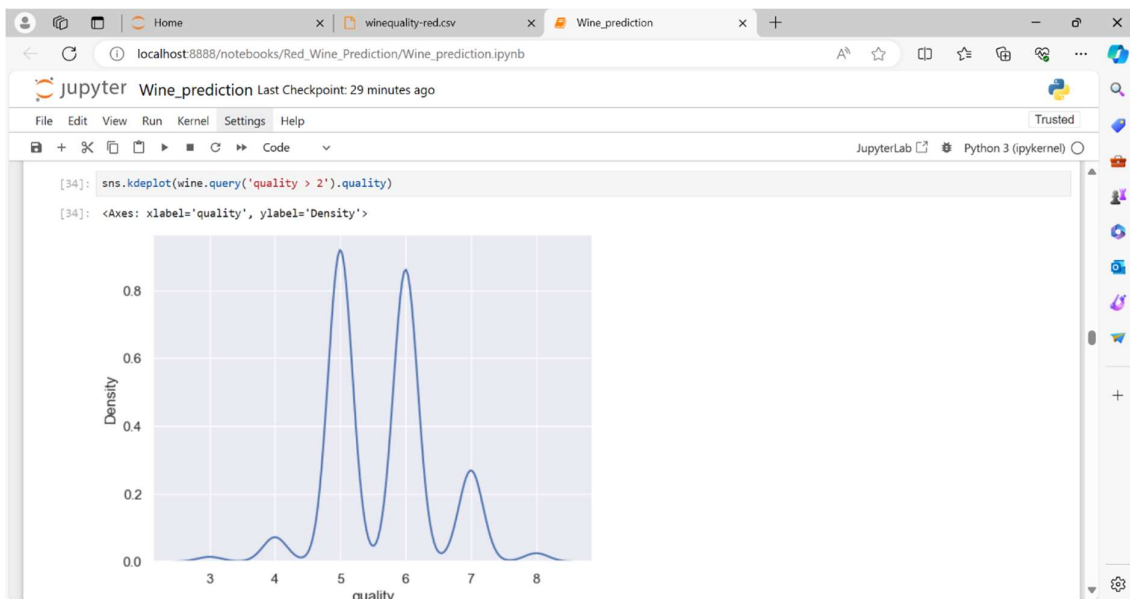

```
sns.countplot(wine['volatile acidity'])  
plt.show()
```

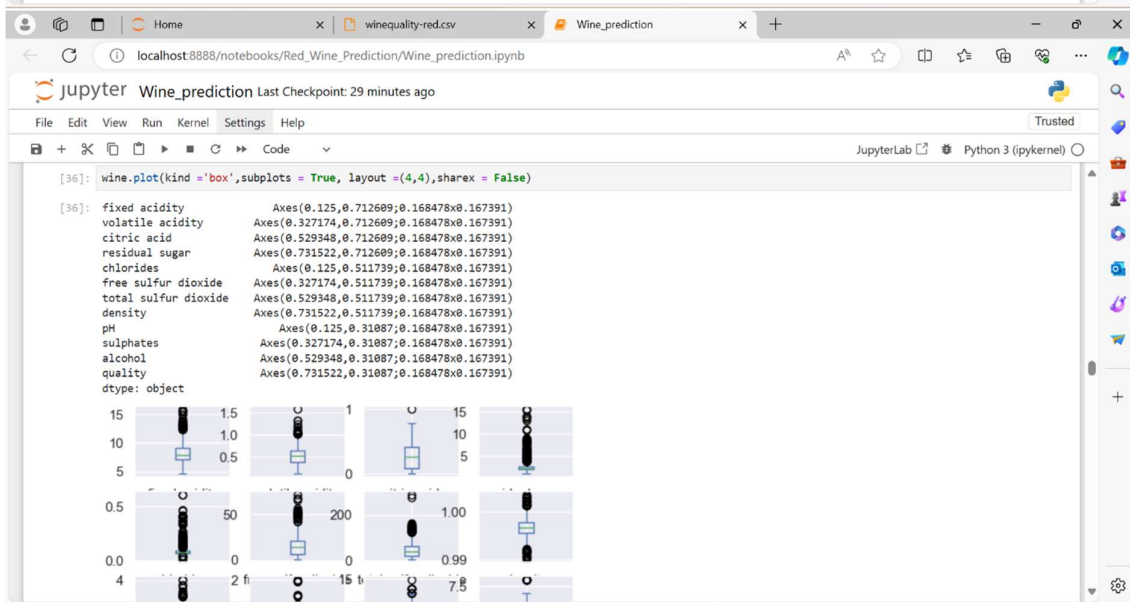
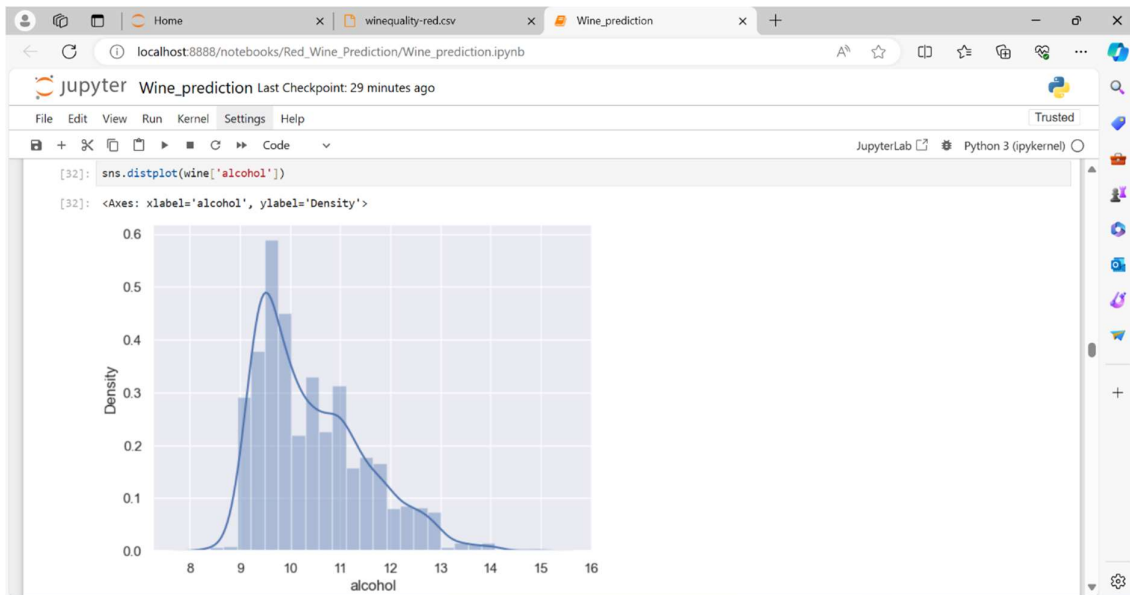


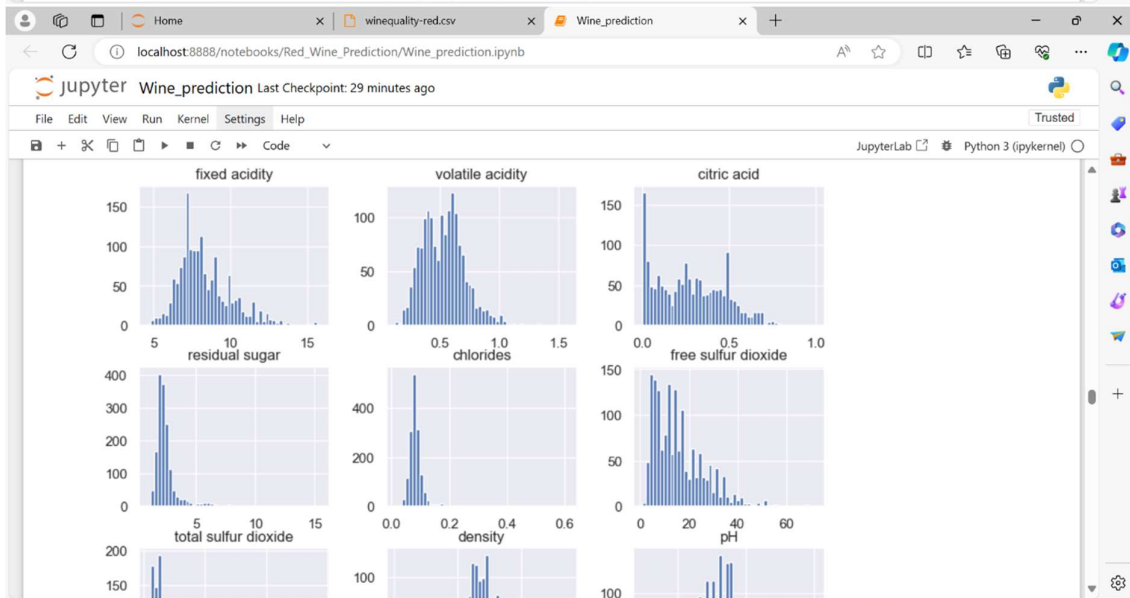
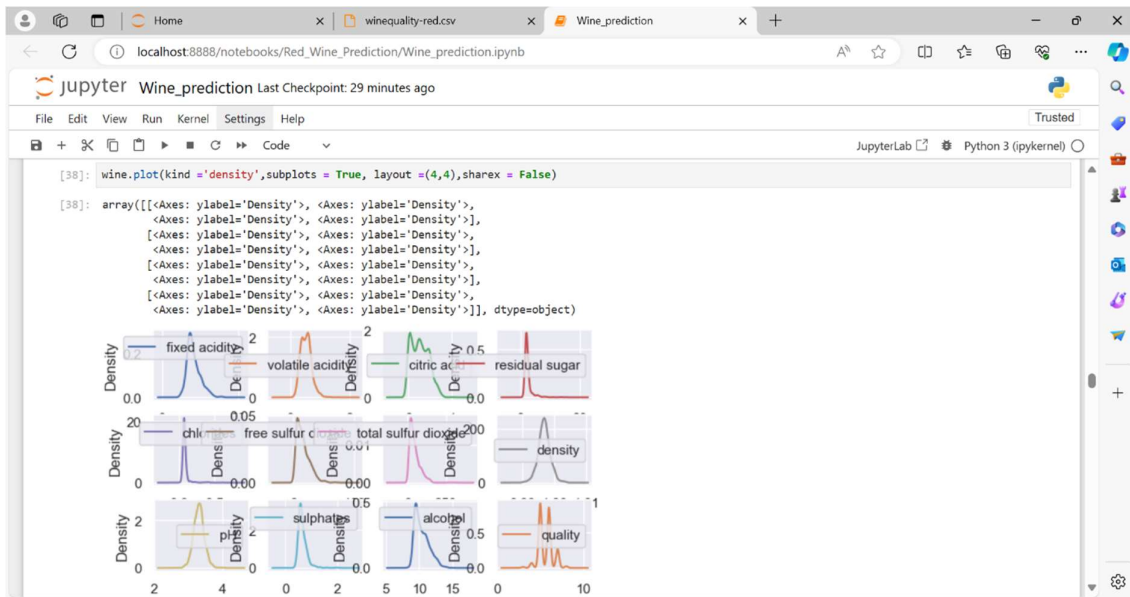
```
sns.countplot(wine['citric acid'])  
plt.show()
```




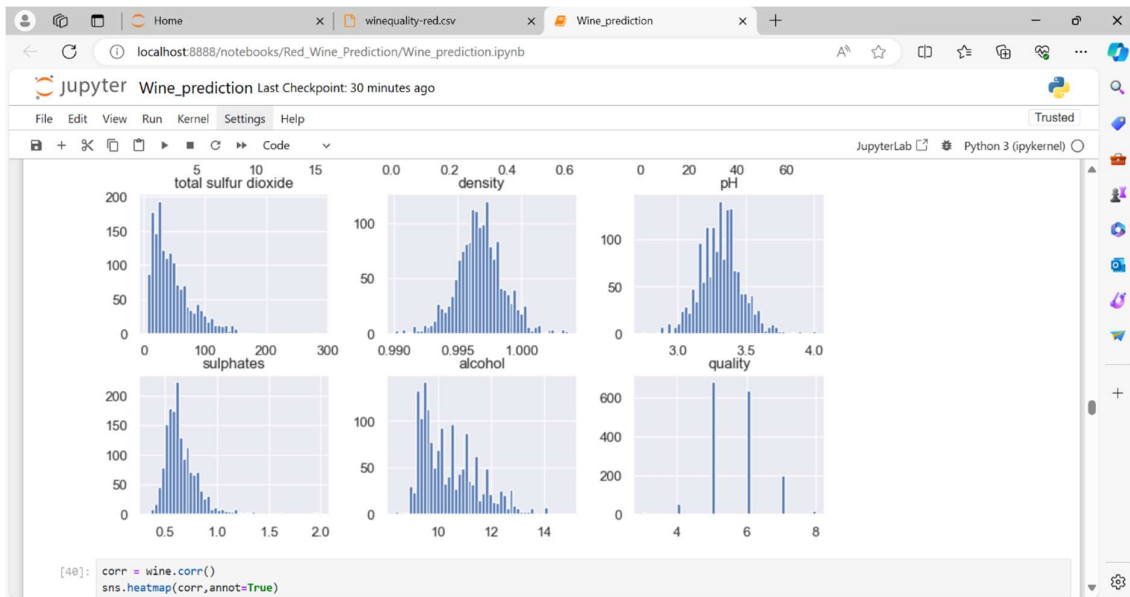
```
sns.countplot(wine['density'])  
plt.show()
```

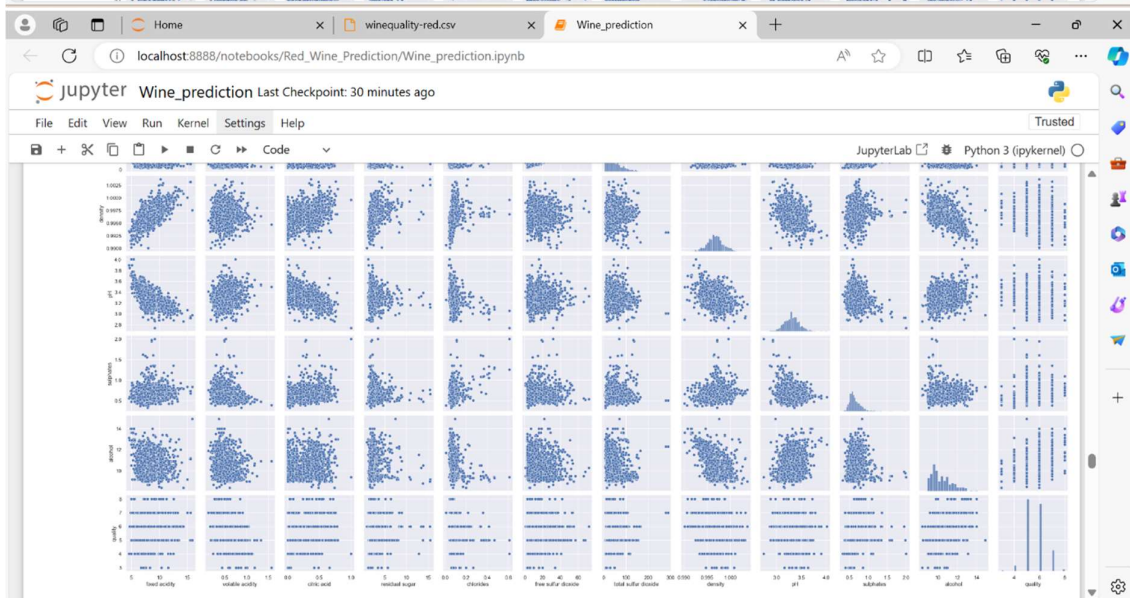
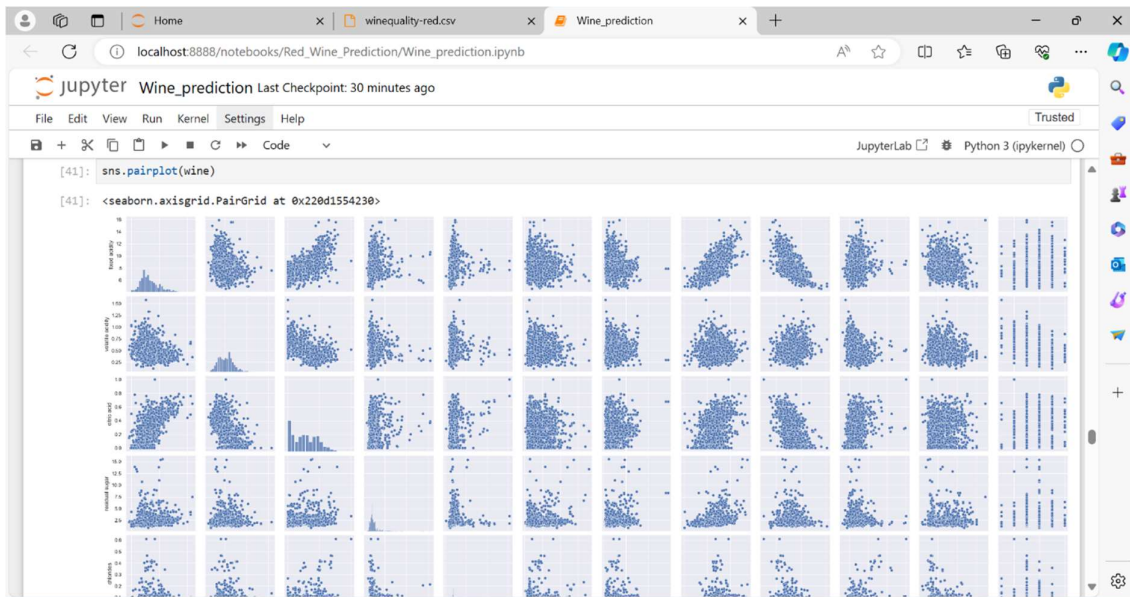


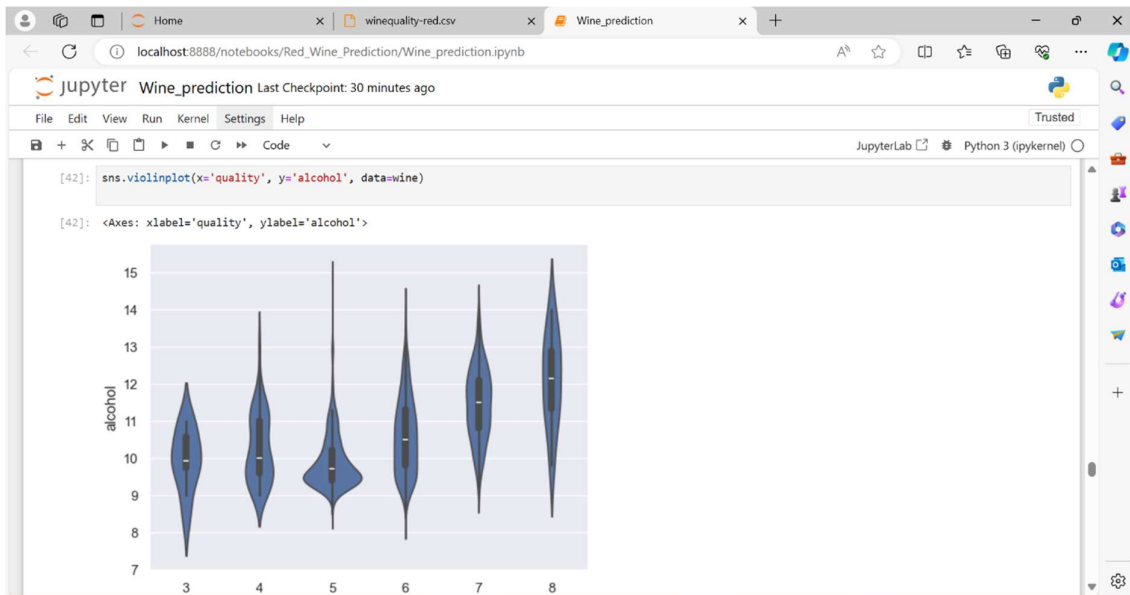












Home | winequality-red.csv | Wine\_prediction | localhost:8888/notebooks/Red\_Wine\_Prediction/Wine\_prediction.ipynb

Jupyter Wine\_prediction Last Checkpoint: 30 minutes ago

File Edit View Run Kernel Settings Help Trusted

JupyterLab Python 3 (ipykernel)

```
[43]: wine['goodquality'] = [1 if x >= 7 else 0 for x in wine['quality']] # Separate feature variables and target variable
X = wine.drop(['quality', 'goodquality'], axis = 1)
Y = wine['goodquality']
```

```
[44]: wine['goodquality'].value_counts()
```

```
[44]: goodquality
0    1382
1     217
Name: count, dtype: int64
```

```
[45]: X
```

```
[45]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
0	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	9.4
1	7.8	0.880	0.00	2.6	0.098	25.0	67.0	0.99680	3.20	0.68	9.8
2	7.8	0.760	0.04	2.3	0.092	15.0	54.0	0.99700	3.26	0.65	9.8
3	11.2	0.280	0.56	1.9	0.075	17.0	60.0	0.99800	3.16	0.58	9.8
4	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	9.4
...	...	...	...	...	...	...	...	...	...	...	...
1504	6.7	0.600	0.08	2.0	0.080	22.0	44.0	0.99690	3.45	0.58	10.5

Home | winequality-red.csv | Wine\_prediction | localhost:8888/notebooks/Red\_Wine\_Prediction/Wine\_prediction.ipynb

Jupyter Wine\_prediction Last Checkpoint: 31 minutes ago

File Edit View Run Kernel Settings Help Trusted

JupyterLab Python 3 (ipykernel)

```
[45]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
0	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	9.4
1	7.8	0.880	0.00	2.6	0.098	25.0	67.0	0.99680	3.20	0.68	9.8
2	7.8	0.760	0.04	2.3	0.092	15.0	54.0	0.99700	3.26	0.65	9.8
3	11.2	0.280	0.56	1.9	0.075	17.0	60.0	0.99800	3.16	0.58	9.8
4	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	9.4
...	...	...	...	...	...	...	...	...	...	...	...
1594	6.2	0.600	0.08	2.0	0.090	32.0	44.0	0.99490	3.45	0.58	10.5
1595	5.9	0.550	0.10	2.2	0.062	39.0	51.0	0.99512	3.52	0.76	11.2
1596	6.3	0.510	0.13	2.3	0.076	29.0	40.0	0.99574	3.42	0.75	11.0
1597	5.9	0.645	0.12	2.0	0.075	32.0	44.0	0.99547	3.57	0.71	10.2
1598	6.0	0.310	0.47	3.6	0.067	18.0	42.0	0.99549	3.39	0.66	11.0

1599 rows x 11 columns

```
[46]: print(Y)
```

```
0    0
1    0
2    0
3    0
4    0
...
1594 0
1595 0
1596 0
1597 0
1598 0
Name: goodquality, Length: 1599, dtype: int64
```

Home | winequality-red.csv | Wine\_prediction | localhost:8888/notebooks/Red\_Wine\_Prediction/Wine\_prediction.ipynb

Jupyter Wine\_prediction Last Checkpoint: 31 minutes ago

File Edit View Run Kernel Settings Help Trusted

JupyterLab Python 3 (ipykernel)

```
[46]: print(Y)
```

```
0    0
1    0
2    0
3    0
4    0
...
1594 0
1595 0
1596 0
1597 0
1598 0
Name: goodquality, Length: 1599, dtype: int64
```

```
[47]: from sklearn.linear_model import LogisticRegression
model = LogisticRegression()

from sklearn.ensemble import ExtraTreesClassifier
classifiern = ExtraTreesClassifier()
classifiern.fit(X,Y)
score = classifiern.feature_importances_
print(score)
```

```
[0.07811046 0.09673209 0.09970649 0.07382695 0.06799365 0.06982716
 0.08338906 0.08286986 0.06508271 0.11441191 0.16884965]
```

```
[48]: from sklearn.model_selection import train_test_split
X_train, X_test, Y_train, Y_test = train_test_split(X,Y,test_size=0.3,random_state=7)
```

Home x winequality-red.csv x Wine\_prediction x +

localhost:8888/notebooks/Red\_Wine\_Prediction/Wine\_prediction.ipynb

Jupyter Wine\_prediction Last Checkpoint: 31 minutes ago

File Edit View Run Kernel Settings Help Trusted

JupyterLab Python 3 (ipykernel)

```
[49]: from sklearn.linear_model import LogisticRegression
      model = LogisticRegression()
      model.fit(X_train,Y_train)
      Y_pred = model.predict(X_test)

      from sklearn.metrics import accuracy_score,confusion_matrix
      print("Accuracy Score:",accuracy_score(Y_test,Y_pred))

      Accuracy Score: 0.86875

[50]: confusion_mat = confusion_matrix(Y_test,Y_pred)
      print(confusion_mat)

      [[399 18]
       [ 45 18]]

[51]: from sklearn.neighbors import KNeighborsClassifier
      model = KNeighborsClassifier(n_neighbors=3)
      model.fit(X_train,Y_train)
      y_pred = model.predict(X_test)

      from sklearn.metrics import accuracy_score
      print("Accuracy Score:",accuracy_score(Y_test,y_pred))

      Accuracy Score: 0.8729166666666667

[52]: from sklearn.svm import SVC
      model = SVC()
      model.fit(X_train,Y_train)
      pred_y = model.predict(X_test)

      from sklearn.metrics import accuracy_score
      print("Accuracy Score:",accuracy_score(Y_test,pred_y))

      Accuracy Score: 0.86875

[53]: from sklearn.tree import DecisionTreeClassifier
      model = DecisionTreeClassifier(criterion='entropy',random_state=7)
      model.fit(X_train,Y_train)
      y_pred = model.predict(X_test)

      from sklearn.metrics import accuracy_score
      print("Accuracy Score:",accuracy_score(Y_test,y_pred))

      Accuracy Score: 0.8645833333333334

[54]: from sklearn.naive_bayes import GaussianNB
      model3 = GaussianNB()
      model3.fit(X_train,Y_train)
      y_pred3 = model3.predict(X_test)

      from sklearn.metrics import accuracy_score
```

Home | winequality-red.csv | Wine\_prediction | localhost:8888/notebooks/Red\_Wine\_Prediction/Wine\_prediction.ipynb

Jupyter Wine\_prediction Last Checkpoint: 31 minutes ago

File Edit View Run Kernel Settings Help Trusted

JupyterLab Python 3 (ipykernel)

Accuracy Score: 0.8333333333333334

```
[55]: from sklearn.ensemble import RandomForestClassifier
model2 = RandomForestClassifier(random_state=1)
model2.fit(X_train, Y_train)
y_pred2 = model2.predict(X_test)

from sklearn.metrics import accuracy_score
print("Accuracy Score:", accuracy_score(Y_test, y_pred2))
```

Accuracy Score: 0.89375

```
[56]: import xgboost as xgb
model5 = xgb.XGBClassifier(random_state=1)
model5.fit(X_train, Y_train)
y_pred5 = model5.predict(X_test)

from sklearn.metrics import accuracy_score
print("Accuracy Score:", accuracy_score(Y_test, y_pred5))
```

Accuracy Score: 0.8895833333333333

```
[58]: results = pd.DataFrame({
    'Model': ['Logistic Regression', 'KNN', 'SVC', 'Decision Tree', 'GaussianNB', 'Random Forest', 'Xgboost'],
    'Score': [0.868, 0.872, 0.868, 0.864, 0.833, 0.893, 0.889]})

result_df = results.sort_values(by='Score', ascending=False)
```

Accuracy Score: 0.8895833333333333

```
[58]: results = pd.DataFrame({
    'Model': ['Logistic Regression', 'KNN', 'SVC', 'Decision Tree', 'GaussianNB', 'Random Forest', 'Xgboost'],
    'Score': [0.868, 0.872, 0.868, 0.864, 0.833, 0.893, 0.889]})

result_df = results.sort_values(by='Score', ascending=False)
result_df = result_df.set_index('Score')
result_df
```

Score	Model
0.893	Random Forest
0.889	Xgboost
0.872	KNN
0.868	Logistic Regression
0.868	SVC
0.864	Decision Tree
0.833	GaussianNB